

Project Topic -2

Reinforcement Learning using WoLF Algorithm

1. Problem Statement:

Two robots are operating in a factory to bring metal bars to two different production halls. The metal bars are dispensed in one place, only one bar can be picked up at a time, and each robot can only carry one bar at a time. Once a metal bar is picked up, a new one will appear at the dispenser with probability 0.5 every time step (every action taken corresponds to one-time step). Each robot has 3 action choices. It can either try to pick up a metal bar, deliver it to the production hall, or wait. If it tries to pick up a metal bar, it will succeed with probability 0.5 (due to imprecisions in its programming) if there is a metal bar available and fail if there is none available. If it tries to deliver a metal bar to the production hall, it will succeed with probability 1 if it is holding a metal bar and fail otherwise. If it decides to wait it will stay in place. If both robots try to pick up a metal bar at the same time, they will both fail. Each robot receives a payoff of 4 if it successfully delivers a metal bar to the production hall and incurs a cost of 1 if it tries to pick up a metal bar or if it tries to deliver one to the production hall (reflecting the energy it uses up). The wait action does not incur a cost.

2. Project Goal:

The goal of this project is to implement the WoLF learning algorithm on the above problem statement. I have applied the WoLF Policy Hill Climbing (WoLF-PHC) algorithm to the problem statement.

3. Introduction to WoLF-PHC:

There are 2 desirable properties of any Multiagent Algorithm:

- **Rationality-** If the other players' policies converge to stationary policies then the learning algorithm will converge to a policy that is a best-response to their policies.
- **Convergence-** The learner will necessarily converge to a stationary policy. This property will usually be conditioned on the other agents using an algorithm from some class of learning algorithms.

With WoLF-PHC, we are able to achieve both the above-mentioned properties. The policy matrix of each agent is converged to an optimal policy matrix and that optimal policy matrix gives us the Nash Equilibrium of each state.

WoLF-PHC Algorithm: -

Given below is the policy hill climbing algorithm for player i . The WoLF-PHC algorithm is the extension of the PHC algorithm. The (a) and (b) step are same. Rest other steps are updated. Please refer to the WoLF-PHC figure below.

1. Let α and δ be learning rates. Initialize,

$$Q(s, a) \leftarrow 0, \quad \pi(s, a) \leftarrow \frac{1}{|\mathcal{A}_i|}.$$

2. Repeat,

(a) From state s select action a with probability $\pi(s, a)$ with some exploration.

(b) Observing reward r and next state s' ,

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha \left(r + \gamma \max_{a'} Q(s', a') \right).$$

(c) Update $\pi(s, a)$ and constrain it to a legal probability distribution,

$$\pi(s, a) \leftarrow \pi(s, a) + \begin{cases} \delta & \text{if } a = \operatorname{argmax}_{a'} Q(s, a') \\ \frac{-\delta}{|\mathcal{A}_i| - 1} & \text{otherwise} \end{cases}.$$

Given below is the WoLF policy hill-climbing algorithm for player.

1. Let $\alpha, \delta_l > \delta_w$ be learning rates. Initialize,

$$Q(s, a) \leftarrow 0, \quad \pi(s, a) \leftarrow \frac{1}{|\mathcal{A}_i|}, \quad C(s) \leftarrow 0.$$

2. Repeat,

(a,b) Same as PHC in Table 1

(c) Update estimate of average policy, $\bar{\pi}$,

$$\begin{aligned} C(s) &\leftarrow C(s) + 1 \\ \forall a' \in \mathcal{A}_i \quad \bar{\pi}(s, a') &\leftarrow \bar{\pi}(s, a') + \frac{1}{C(s)} (\pi(s, a') - \bar{\pi}(s, a')). \end{aligned}$$

(d) Update $\pi(s, a)$ and constrain it to a legal probability distribution,

$$\pi(s, a) \leftarrow \pi(s, a) + \begin{cases} \delta & \text{if } a = \operatorname{argmax}_{a'} Q(s, a') \\ \frac{-\delta}{|\mathcal{A}_i| - 1} & \text{otherwise} \end{cases},$$

where,

$$\delta = \begin{cases} \delta_w & \text{if } \sum_a \pi(s, a) Q(s, a) > \sum_a \bar{\pi}(s, a) Q(s, a) \\ \delta_l & \text{otherwise} \end{cases}.$$

4. Results of my project: -

I followed the algorithm as stated above. In the project we have 8 states and 3 actions. I have done this project in Python 3.

- Before I start explaining what I did in the project, I want to tell that the system has 8 states and each agent (Robot_1 and Robot_2 in our case) has 3 actions

		Dispenser	R1	R2
States	S1	0	0	0
	S2	0	0	1
	S3	0	1	0
	S4	0	1	1
	S5	1	0	0
	S6	1	0	1
	S7	1	1	0
	S8	1	1	1

Actions	Robot Waits	Robot Picks up	Robot Delivers
---------	-------------	----------------	----------------

- I have initialized the value of alpha to be 0.01, discount factor to be 0.9, delta win to be 0.2 and delta loose to be 0.8
- I initialized the $Q(s,a)$ for both the robots. This returned a numpy array of dimension (8 X 3) with zero values as shown in screenshot below.

```
[[ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]
```

Figure 1. $Q(s,a) \rightarrow 0$

- Then, I initialized the $\pi(s,a)$ matrix with the values 0 because as per the algorithm, it should be initialized with the values equal to $1/|A|$. Since the cardinality of the action set in my project is 3, I have initialized it with 0.33 as show in the figure below.

```

[[ 0.33333333  0.33333333  0.33333333]
 [ 0.33333333  0.33333333  0.33333333]
 [ 0.33333333  0.33333333  0.33333333]
 [ 0.33333333  0.33333333  0.33333333]
 [ 0.33333333  0.33333333  0.33333333]
 [ 0.33333333  0.33333333  0.33333333]
 [ 0.33333333  0.33333333  0.33333333]
 [ 0.33333333  0.33333333  0.33333333]]

```

Figure 2- $\pi(s,a) = 1/|A|$

- Then I initialized the counter constant $C(s)$ to 0 as suggested in the algorithm. The goal of the counter is to average the policy in step (c) of the algorithm.
- The next step is to update the q values of both the agents using the equation given in the (b) part of the Policy Hill Climbing algorithms. I calculated and updated the q values using the same equation. Please refer to the screenshots below for the final values of both the robots.

```

Final Q values for Robot 1

[[ 1.10671098  0.2419109  0.2419109 ]
 [ 2.42632421  1.56152414  1.56152414]
 [ 1.73721321  0.87241313  4.33161343]
 [ 0.41760489 -0.44719519  3.01200511]
 [ 1.10673502  0.24193494  0.24193494]
 [ 2.42640721  1.56160714  1.56160714]
 [ 1.73745466  0.87265459  4.33185489]
 [ 0.41817732 -0.44662276  3.01257754]]

```

Final Q values for Robot 2

```
[ [ 1.94728548  1.0824854  1.0824854 ]  
 [ 0.89663871  0.03183864  3.49103893]  
 [ 1.94729039  1.08249031  1.08249031]  
 [ 0.89665784  0.03185777  3.49105807]  
 [ 1.94735425  1.08255418  1.08255418]  
 [ 0.89683543  0.03203536  3.49123565]  
 [ 1.9477491   1.08294902  1.08294902]  
 [ 0.89749353  0.03269345  3.49189375] ]
```

- The next step is to update the estimate of the average policy using the equation mentioned step (c) of WoLF algorithm. I updated using the same equation.

Final Policy matrix for Robot 1

```
[ [-52.01916667 -52.04666667 -52.04666667]  
 [-76.52166667 -76.71416667 -76.68666667]  
 [-89.66666667 -89.66666667 -89.39166667]  
 [-89.66666667 -89.66666667 -89.33666667]  
 [-52.01916667 -52.04666667 -52.04666667]  
 [-76.52166667 -76.71416667 -76.68666667]  
 [-89.66666667 -89.66666667 -89.39166667]  
 [-89.66666667 -89.66666667 -89.33666667] ]
```

Final Policy matrix for Robot 2

```
[ [-74.54166667 -75.06416667 -75.00916667]  
 [-89.66666667 -89.66666667 -89.33666667]  
 [-74.54166667 -75.06416667 -75.00916667]  
 [-89.66666667 -89.66666667 -89.33666667]  
 [-74.54166667 -75.06416667 -75.00916667]  
 [-89.66666667 -89.66666667 -89.33666667]  
 [-74.54166667 -75.06416667 -75.00916667]  
 [-89.66666667 -89.66666667 -89.33666667] ]
```

- The final step is to update the policy matrix $\pi(s,a)$ in order to get the converged policy matrix. This final policy matrix for the both the agents leads to the Nash Equilibrium of all the states. I was successfully able to achieve the correct Nash Equilibrium of all the states. Please refer the screenshots below.

Nash Equilibrium for State 0
wait , wait

Nash Equilibrium for State 1
wait , deliver

Nash Equilibrium for State 2
deliver , wait

Nash Equilibrium for State 3
deliver , deliver

Nash Equilibrium for State 4
wait , wait

Nash Equilibrium for State 5
wait , deliver

Nash Equilibrium for State 6
deliver , wait

Nash Equilibrium for State 7
deliver , deliver

- Later on, after reaching the convergence of the optimal policy matrix (in around 2000 iterations per agent), I was able to get the Nash Equilibrium for all the states.

5. References: -

- a. <http://www.cs.cmu.edu/~mmv/papers/01ijcai-mike.pdf>
- b. <http://www.masfoundations.org/mas.pdf>